

CS 452 – Kernel 3

February 6, 2017

Zhengkun Chen

Student ID: 20557054

Xiwen Liu

Student ID: 20468618

Operating Instructions

Full pathname to executable: /u/cs452/tftp/ARM/z283chen/kernel.elf

To load and run the program in the Redboot terminal, run the following command:

```
> load -b 0x00218000 -h 10.15.167.5 "ARM/z283chen/kernel.elf" ; go
```

The executable can also be created by going into the directory “kernel” and calling make:

```
> cd kernel
> make
```

And then copy the elf to wherever you want to run it.

Kernel Primitives

All kernel primitives were implemented as specified in the kernel 3 assignment documents.

Kernel Structure

Context Switch

- Software Interrupts
 - During initialization, the label for the software interrupt handler is loaded into address 0x28.
 - Context switching was implemented with inline assembly.
 - Entering kernel: The kernel saves user's registers onto user stack, and retrieve registers from its own stack.
 - Exiting kernel: The kernel saves its registers onto the kernel stack, retrieve registers for user program and load link register into pc and then hands over the control to user program.
- Hardware Interrupts
 - During initialization, the label for the IRQ handler is loaded into address 0x38. The SWI handler is a subset of the IRQ handler.
 - Entering kernel: On entering the kernel, the IRQ handler will first save user state, including the sp. Then it transfers the spsr and lr of IRQ mode into SVC mode. After this, it runs the same instructions as SWI.
 - Exiting kernel: If the kernel is exiting from an interrupt, it will load back all of the user registers it initially saved, as well as the user lr. Then the rest of the instructions are the same as SWI.

User Tasks

Algorithms

- Creating task descriptors
 - Task descriptors are located at a reserved memory address that is slightly under kernel stack. A task descriptor contains information for task to run, pause, and resume, including link register (when entering kernel), return value of a syscall, and much more.
- Deleting tasks
 - When a task calls Exit(), it is removed from the priority queue.
 - Task descriptors are not removed from memory. A maximum of 50 tasks can be created.

Data Structures

- User task stack: This stack starts at 0x01200000, and each user task is allocated a stack size of 0x00015000 bytes. The stack grows downward as per GCC convention.

- Task Descriptors: An array of task descriptors begins at 0x01400000, which is a space reserved only for task descriptors. Using a global variable could be an alternative but it will be more work to access them via assembly code. A maximum of 50 tasks and task descriptors can be created.
 - A copy of the active task's task descriptor is maintained at address 0x01300000. This copy helps simplify retrieving information needed about the active task, such as its task id and priority.

Scheduling

Algorithms

- The scheduler chooses the task with the highest priority. The scheduling functions require the priority of the task as a parameter. Adding a new task to the priority queues, removing a task, scheduling a new task, and rescheduling a task each take $O(1)$ time.
- Round robin scheduling is used for adding new tasks and rescheduling tasks, so the task is always added to the back of the queue.

Data Structures

- The priority queue structure contains a circular queue of task id's, and the indices of the first and last values in the task id array. The size of the task id array is set to the maximum number of tasks.
- An array of priority queue structures is allocated beginning at address 0x01600000, and grows upward. Each element in this array is a priority queue for a different priority. Priorities can be dynamically added or removed by changing the size of this array. There are 32 priorities available.
- Since the scheduling functions take in a priority as a parameter, the array allows the correct queue to be selected in $O(1)$ time. The first and last indices in the priority queue structure allow functions to add and remove tasks in $O(1)$ time as well.

Message Passing

Algorithms

- Send before Receive: send adds the current task to receiver task's sendQ; blocks itself. And then receiver task will extract the info from its sendQ. Copy the content of the message from sender's space to receiver's space (done by kernel). And then receiver should call reply accordingly to send back the reply to the sender, and unblocks the sender.
- Receive before Send: the receiver first blocks itself, we implemented this by setting the state of current task to blocked and calling a pass. And once a sender is detected, it is resumed.

Data Structures

- The sendQ structure is a circular queue, with pointers to the first and last elements in the queue.

Interrupt Handling

- AwaitEvent(): It sets the current task to EVENT_BLOCKED. During the IRQ handling, the corresponding task (which is usually a notifier) that is waiting on the interrupt will be unblocked and resumed to process the data received from that interrupt.

Nameserver

Algorithms

- The nameserver performs linear searching for WhoIs() and linear insertion for RegisterAs(). The worst case time is $O(n)$, where n is the number of names that have been inserted. The maximum number of names that can be inserted is 50, so the worst case time is short for both operations.
- Our implementation of strcpy, which is identical to its implementation in string.h, was used to copy the given name to the nameserver structure.

Data Structures

- The nameserver begins at address 0x018000000, and grows up. It is an array of nameserver structures, which contain the fields name and tid. Name a C-style string and tid is an integer.
- The array is not circular, and names cannot be removed, so a maximum of 50 names can be registered.

Clock server

Algorithms

- The clock server functions Time, Delay, and DelayUntil use the message passing primitives Send, Receive, and Reply to block the calling task until it receives the requested data. The clock server calls receive first and is receive blocked until a client calls one of Time, Delay, or DelayUntil, which will then call Send and so on.
- Delay is a wrapper for DelayUntil. It adds the current time to the input time, and then calls DelayUntil with the new time. DelayTime then causes the server to add the task to the delay queue. Calling Delay or DelayUntil blocks the task until the calculated system time has been reached.
- New delayed tasks are added to the delay queue with insertion sort, so the task with the shortest delay time is always at the front of the queue. This runs in $O(n)$. Removing a task from the queue also takes $O(n)$, since it is not a circular array. We intend to replace these structures and algorithms with more efficient ones when the load on the program increases.

Data Structures

- For delaying tasks, the clock servers uses an array of structures. This structure contains the fields tid and delay time. The delay time is the system time of when the task should be unblocked.
- The clock server uses Timer3 to keep time.

Idle Task

- The idle task is set at the lowest priority. When it first runs, it initiates Timer1, and then enters a while loop. The idle task keeps track of how many ticks has passed while it is running. The idle usage is calculated by dividing the amount of ticks counted by the idle task over the amount of ticks counted by the program.

Source Code

Git repository: <https://git.uwaterloo.ca/x272liu/cs452>

Sha1 of the commit: f8f2c18e50e8118dd939f5c87a749dd89bbdd55e

All files in the repository are part of this submission.

Program Output

```
Task 6 created at: 0
Task 7 created at: 0
Task 8 created at: 0
Task 9 created at: 0
Tid 6: delay time is 10 number of delays is 1      idle usage: 90%
Tid 6: delay time is 10 number of delays is 2      idle usage: 91%
Tid 7: delay time is 23 number of delays is 1      idle usage: 89%
Tid 6: delay time is 10 number of delays is 3      idle usage: 89%
Tid 8: delay time is 33 number of delays is 1      idle usage: 88%
Tid 6: delay time is 10 number of delays is 4      idle usage: 88%
Tid 7: delay time is 23 number of delays is 2      idle usage: 88%
Tid 6: delay time is 10 number of delays is 5      idle usage: 88%
Tid 6: delay time is 10 number of delays is 6      idle usage: 88%
Tid 8: delay time is 33 number of delays is 2      idle usage: 88%
Tid 7: delay time is 23 number of delays is 3      idle usage: 88%
Tid 6: delay time is 10 number of delays is 7      idle usage: 87%
Tid 9: delay time is 71 number of delays is 1      idle usage: 86%
Tid 6: delay time is 10 number of delays is 8      idle usage: 87%
Tid 6: delay time is 10 number of delays is 9      idle usage: 87%
Tid 7: delay time is 23 number of delays is 4      idle usage: 87%
Tid 8: delay time is 33 number of delays is 3      idle usage: 87%
Tid 6: delay time is 10 number of delays is 10     idle usage: 86%
Tid 6: delay time is 10 number of delays is 11     idle usage: 87%
Tid 7: delay time is 23 number of delays is 5      idle usage: 87%
Tid 6: delay time is 10 number of delays is 12     idle usage: 87%
Tid 6: delay time is 10 number of delays is 13     idle usage: 87%
Tid 8: delay time is 33 number of delays is 4      idle usage: 87%
Tid 7: delay time is 23 number of delays is 6      idle usage: 87%
Tid 6: delay time is 10 number of delays is 14     idle usage: 86%
Tid 9: delay time is 71 number of delays is 2      idle usage: 86%
Tid 6: delay time is 10 number of delays is 15     idle usage: 86%
Tid 6: delay time is 10 number of delays is 16     idle usage: 87%
Tid 7: delay time is 23 number of delays is 7      idle usage: 86%
Tid 8: delay time is 33 number of delays is 5      idle usage: 86%
Tid 6: delay time is 10 number of delays is 17     idle usage: 86%
Tid 6: delay time is 10 number of delays is 18     idle usage: 86%
Tid 7: delay time is 23 number of delays is 8      idle usage: 86%
Tid 6: delay time is 10 number of delays is 19     idle usage: 86%
Tid 8: delay time is 33 number of delays is 6      idle usage: 87%
Task 8 exiting at: 198
Tid 6: delay time is 10 number of delays is 20     idle usage: 86%
Task 6 exiting at: 200
Tid 7: delay time is 23 number of delays is 9      idle usage: 86%
Task 7 exiting at: 207
Tid 9: delay time is 71 number of delays is 3      idle usage: 86%
Task 9 exiting at: 213      idle usage: 86%
```

Explanation of output

The order of printed statements should be in the order of the amount of each delay time multiplied by the number of delays, and as is the exit time. E.g. “Tid 7: delay time is 23 number of delays is 1” is before “Tid 6: delay time is 10 number of delays is 3” since 23 times 1 is less than 10 times 3. The finishing order should correspond to the total delay time.

| Task | Priority | Delay Time | Number of Delays | Total Delay Time |
|------|----------|------------|------------------|------------------|
| 6 | 3 | 10 | 20 | 200 |
| 7 | 4 | 23 | 9 | 207 |
| 8 | 5 | 33 | 6 | 198 |
| 9 | 6 | 71 | 3 | 213 |

The percentage printed is the portion of time that idle task is running compared to the total uptime of the entire program. In our case, since most tasks are just constantly calling Delay(), most of the CPU time should be spent on idle task as expected.